

## Lecture 3: Low-rank matrix approximation



School of Mathematical Sciences, Xiamen University

## 1. Low-rank matrix approximation problem

- Suppose  $\mathbf{B}$  is a rectangular matrix, accessible via matvecs  $\mathbf{x} \mapsto \mathbf{B}\mathbf{x}$  and  $\mathbf{y} \mapsto \mathbf{B}^*\mathbf{y}$ . The task is to produce a low-rank approximation of  $\mathbf{B}$  that is competitive with a best approximation of similar rank.
- The best rank- $k$  approximation is unique if and only if  $\sigma_k > \sigma_{k+1}$ :

$$\min_{\text{rank}(\mathbf{M}) \leq k} \|\mathbf{B} - \mathbf{M}\|_{\text{F}}^2 = \|\mathbf{B} - \mathbf{U}_k \mathbf{U}_k^* \mathbf{B}\|_{\text{F}}^2 = \sum_{i > k} \sigma_i^2,$$

where  $\mathbf{U}_k$  is the matrix consisting of the leading  $k$  left singular vectors. Cost:  $\mathcal{O}(mnp)$  where  $p = \min(m, n)$ .

- Let  $s = k + l$  for a small natural number  $l$ . For a tolerance  $\varepsilon > 0$ , we seek a rank- $s$  approximation  $\hat{\mathbf{B}}_s$  that competes with the best rank- $k$  approximation:

$$\|\mathbf{B} - \hat{\mathbf{B}}_s\|_{\text{F}}^2 \leq (1 + \varepsilon) \|\mathbf{B} - \mathbf{U}_k \mathbf{U}_k^* \mathbf{B}\|_{\text{F}}^2 = (1 + \varepsilon) \sum_{i > k} \sigma_i^2.$$

## 1.1 Randomized SVD: Intuition

- Draw a standard normal test vector  $\mathbf{w} \in \mathbb{R}^n$ . we have

$$\mathbf{B}\mathbf{w} = \sum_{i=1}^p \sigma_i \mathbf{u}_i (\mathbf{v}_i^* \mathbf{w}) := \sum_{i=1}^p \sigma_i \mathbf{u}_i w_i.$$

The component  $w_i := \mathbf{v}_i^* \mathbf{w}$  of the random vector along the  $i$ th right singular vector follows a standard normal distribution, and the components  $(w_i, i = 1, \dots, p)$  compose an independent family.

- On average,  $\mathbb{E}(w_i^2) = 1$ . Therefore, the image  $\mathbf{B}\mathbf{w}$  tends to align with the left singular vectors associated with large singular values.
- By repeating this process with a statistically independent family  $(\mathbf{w}^{(j)} : j = 1, \dots, s)$  of random test vectors, we can obtain a family  $(\mathbf{B}\mathbf{w}^{(j)} : j = 1, \dots, s)$  of vectors whose span contains most of  $\text{range}(\mathbf{U}_k)$ . The number  $s = k + l$  of test vectors needs to be a bit larger than the target rank  $k$  to obtain coverage of the subspace with high probability.

## 1.2 Randomized SVD: Algorithm (cost $\mathcal{O}(smn)$ )

- For a rank parameter  $s$ , we draw a random test matrix:

$$\mathbf{\Omega} = [\mathbf{w}^{(1)} \quad \dots \quad \mathbf{w}^{(s)}] \quad \text{where} \quad \mathbf{w}^{(j)} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_n) \quad \text{i.i.d.}$$

- We obtain  $\mathbf{Y} := \mathbf{B}\mathbf{\Omega}$ . The orthogonal projector  $\mathbf{P}_{\mathbf{Y}}$  onto  $\text{range}(\mathbf{Y})$  serves as a proxy for the ideal projector  $\mathbf{U}_k \mathbf{U}_k^*$ . Computationally,

$$\mathbf{P}_{\mathbf{Y}} := \mathbf{Q}\mathbf{Q}^* \quad \text{where} \quad \mathbf{Q} := \text{orth}(\mathbf{Y}).$$

The function `orth` returns an orthonormal basis and costs  $\mathcal{O}(s^2m)$ .

- Finally, we report the approximation  $\hat{\mathbf{B}}_s$  in factored form:

$$\hat{\mathbf{B}}_s := \mathbf{P}_{\mathbf{Y}}\mathbf{B} = \mathbf{Q}(\mathbf{Q}^*\mathbf{B}).$$

- If desired, we can report the SVD of the approximation after a small amount of additional work ( $\mathcal{O}(s^2n)$ ):

$$\hat{\mathbf{B}}_s = (\mathbf{Q}\hat{\mathbf{U}}_0)\hat{\mathbf{\Sigma}}\hat{\mathbf{V}}^* \quad \text{where} \quad (\hat{\mathbf{U}}_0, \hat{\mathbf{\Sigma}}, \hat{\mathbf{V}}) = \text{svd}(\mathbf{Q}^*\mathbf{B}).$$

---

**Algorithm:** Randomized SVD.

---

$$\mathbf{\Omega} = \text{randn}(n, s)$$

$$\mathbf{Y} = \mathbf{B}\mathbf{\Omega}$$

$$\mathbf{Q} = \text{orth}(\mathbf{Y})$$

$$\mathbf{C} = \mathbf{Q}^* \mathbf{B}$$

$$(\hat{\mathbf{U}}_0, \hat{\mathbf{\Sigma}}, \hat{\mathbf{V}}) = \text{svd}(\mathbf{C})$$

$$\hat{\mathbf{U}} = \mathbf{Q} \hat{\mathbf{U}}_0$$

---

### Theorem 1

Consider a matrix  $\mathbf{B} \in \mathbb{R}^{m \times n}$ , and fix the target rank  $k \leq p$ . When  $s \geq k + 2$ , the randomized SVD method produces a random rank- $s$  approximation  $\hat{\mathbf{B}}_s$  that satisfies

$$\mathbb{E} \|\mathbf{B} - \hat{\mathbf{B}}_s\|_{\text{F}}^2 \leq \left(1 + \frac{k}{s - k - 1}\right) \sum_{i > k} \sigma_i^2(\mathbf{B}).$$

Proof. See A. Kireeva and J. A. Tropp, arXiv:2402.17873, 2024. □

## 1.3 Randomized subspace iteration

---

**Algorithm:** Randomized subspace iteration.

---

```
 $\mathbf{X}_0 = \text{randn}(n, s)$   
for  $t = 1, 2, \dots, T$   
     $\mathbf{Q}_t := \text{orth}(\mathbf{B}\mathbf{X}_{t-1})$   
     $\mathbf{X}_t := \mathbf{B}^* \mathbf{Q}_t$   
end  
 $\hat{\mathbf{B}}_s := \mathbf{Q}_T \mathbf{X}_T^*$ 
```

---

- Randomized SVD is the special case of this algorithm with  $T = 1$ .
- Randomized subspace iteration produces approximations

$$\hat{\mathbf{B}}_s = \mathbf{Q}_t(\mathbf{Q}_t^* \mathbf{B}) \quad \text{where} \quad \mathbf{Q}_t = \text{orth}((\mathbf{B}\mathbf{B}^*)^{t-1} \mathbf{B}\mathbf{\Omega}) \quad \text{for } t = 1, 2, \dots$$

- Much as the block power method drives its iterates toward the leading eigenspace, subspace iteration drives  $\text{range}(\mathbf{Q}_t)$  so that it aligns with  $\text{range}(\mathbf{U}_r)$ , the leading left singular subspace of  $\mathbf{B}$ .

## 2. Low-rank spsd approximation from entries

Consider an spsd matrix  $\mathbf{A} \in \mathbb{R}^{n \times n}$  that we access via entry evaluations:  $(j, k) \mapsto a_{jk}$ . The task is to produce a low-rank spsd approximation of  $\mathbf{A}$  using as few as entry evaluations.

### 2.1 Column Nyström approximation

- Given a list  $S \subseteq \{1, 2, 3, \dots, n\}$  of column indices, the column Nyström approximation:

$$\mathbf{A}_{\langle S \rangle} := \mathbf{A}(:, S) \mathbf{A}(S, S)^\dagger \mathbf{A}(S, :).$$

- The column Nyström approximation has several remarkable properties: (1)  $\text{range}(\mathbf{A}_{\langle S \rangle}) = \text{range}(\mathbf{A}(:, S))$ ; (2)  $\mathbf{0} \preceq \mathbf{A}_{\langle S \rangle} \preceq \mathbf{A}$ .
- Our goal is to find a set  $S$  of  $s$  columns that make the error

$$\|\mathbf{A} - \mathbf{A}_{\langle S \rangle}\|_* = \text{tr}(\mathbf{A} - \mathbf{A}_{\langle S \rangle})$$

as small as possible.

## 2.2 Pivoted partial Cholesky

- Set  $\mathbf{A}_0 := \mathbf{A}$  and  $\hat{\mathbf{A}}_0 := \mathbf{0}$ .

At each step  $t = 1, 2, \dots, s$ , select  $i_t \in \{1, 2, \dots, n\}$ , and update

$$\mathbf{A}_t := \mathbf{A}_{t-1} - \frac{\mathbf{A}_{t-1}(:, i_t) \mathbf{A}_{t-1}(i_t, :)}{\mathbf{A}_{t-1}(i_t, i_t)};$$

$$\hat{\mathbf{A}}_t := \hat{\mathbf{A}}_{t-1} + \frac{\mathbf{A}_{t-1}(:, i_t) \mathbf{A}_{t-1}(i_t, :)}{\mathbf{A}_{t-1}(i_t, i_t)}.$$

**Exercise:** Prove the following results: (i)  $\mathbf{A}_t + \hat{\mathbf{A}}_t = \mathbf{A}$ ;

(ii)  $\text{diag}(\mathbf{A}_t) = \text{diag}(\mathbf{A}_{t-1}) - \frac{1}{\mathbf{A}_{t-1}(i_t, i_t)} |\mathbf{A}_{t-1}(:, i_t)|^2.$

### Proposition 2

*Suppose that we apply the pivoted partial Cholesky algorithm to an spsd matrix  $\mathbf{A}$ , and we select  $i_t$  from  $S$  in any order. Then  $\hat{\mathbf{A}}_{|S|} = \mathbf{A}_{\langle S \rangle}$ , where  $|S|$  denotes the number of elements of the set  $S$ .*

The proof is left as an exercise.



## 2.3 Pivoted partial Cholesky: evaluating fewer entries

- Set  $\mathbf{F}_0 := \mathbf{0}$ .

At each step  $t = 1, 2, \dots, s$ , select  $i_t \in \{1, 2, \dots, n\}$  and set

$$\mathbf{c}_t := \mathbf{A}(:, i_t) - \mathbf{F}_{t-1}(\mathbf{F}_{t-1}(i_t, :))^*.$$

Update  $\mathbf{F}_t := [\mathbf{F}_{t-1} \quad \mathbf{c}_t / \sqrt{\mathbf{c}_t(i_t)}]$ .

**Exercise:** Prove that  $\hat{\mathbf{A}}_t = \mathbf{F}_t \mathbf{F}_t^*$  for  $t = 0, 1, 2, \dots, s$ .

## 2.4 Pivot selection rules

- Uniform random pivoting:  $i_t \sim \text{uniform}\{1, 2, \dots, n\}$ .

Assumption: data points represent an i.i.d. sample from a population, so one is just as good as another.

- Greedy pivoting:  $i_t \in \operatorname{argmax}\{\mathbf{A}_{t-1}(i, i) : i = 1, 2, \dots, n\}$ .

Note that  $\mathbf{A}_t(i_t, i_t) = 0$ .

- Importance sampling pivoting:  $\mathbb{P}\{i_t = j\} = \mathbf{A}_{t-1}(j, j) / \operatorname{tr}(\mathbf{A}_{t-1})$ .

Balance between uniform random and greedy.

## 2.5 Randomly pivoted partial Cholesky

---

**Algorithm:** Randomly pivoted partial Cholesky.

---

$\mathbf{F} = \text{zeros}(n, s)$  (Preallocation)  
 $\mathbf{d} = \text{diag}(\mathbf{A})$   
**for**  $t = 1, 2, \dots, s$   
    Sample  $i_t \sim \mathbf{d} / \sum_{j=1}^n \mathbf{d}(j)$   
     $\mathbf{c} = \mathbf{A}(:, i_t) - \mathbf{F}(\mathbf{F}(i_t, :))^*$   
     $\mathbf{F}(:, t) = \mathbf{c} / \sqrt{\mathbf{c}(i_t)}$   
     $\mathbf{d} = \mathbf{d} - |\mathbf{F}(:, t)|^2$   
     $\mathbf{d} = \max\{\mathbf{d}, \mathbf{0}\}$  (Improve numerical stability)  
    Stop when  $\sum_{j=1}^n \mathbf{d}(j) < \eta \cdot \text{tr}(\mathbf{A})$  (Optional)  
**end**

---

- To produce a rank- $s$  approximation, the algorithm only requires  $(s + 1)n - s$  entries of  $\mathbf{A}$ : its diagonal and the  $s$  pivot columns.

- Define the expected residual map:  $\Phi(\mathbf{A}) := \mathbb{E}(\mathbf{A}_1)$ . This function measures the average progress that we make after one step of the algorithm. A quick calculation yields a formula for the expected residual map:

$$\begin{aligned}\Phi(\mathbf{A}) &= \sum_{j=1}^n \left[ \mathbf{A} - \frac{\mathbf{A}(:,j)\mathbf{A}(j,:)}{\mathbf{A}(j,j)} \right] \frac{\mathbf{A}(j,j)}{\text{tr}(\mathbf{A})} \\ &= \mathbf{A} - \frac{1}{\text{tr}(\mathbf{A})} \sum_{j=1}^n \mathbf{A}(:,j)\mathbf{A}(j,:) = \mathbf{A} - \frac{\mathbf{A}^2}{\text{tr}(\mathbf{A})}.\end{aligned}$$

As a result, we have

$$\mathbb{E}(\text{tr}(\mathbf{A}_1)) = \text{tr}(\mathbb{E}(\mathbf{A}_1)) = \left(1 - \frac{\text{tr}(\mathbf{A}^2)}{(\text{tr}(\mathbf{A}))^2}\right) \text{tr}(\mathbf{A}) \leq \frac{n-1}{n} \text{tr}(\mathbf{A}).$$

In each iteration, we decrease the expected trace of the residual on average.

- The best rank- $k$  approximation in  $\|\cdot\|_*$ :

$$\min_{\text{rank}(\mathbf{M}) \leq k} \|\mathbf{A} - \mathbf{M}\|_* = \sum_{j>k} \sigma_j(\mathbf{A}).$$

- Fix a comparison rank  $k$  and a tolerance  $\varepsilon > 0$ . Randomly pivoted Cholesky produces an approximation  $\hat{\mathbf{A}}_s$  that attains the error bound

$$\mathbb{E}(\|\mathbf{A} - \hat{\mathbf{A}}_s\|_*) \leq (1 + \varepsilon) \sum_{j>k} \sigma_j(\mathbf{A})$$

after selecting  $s$  columns where

$$s \geq \frac{k}{\varepsilon} + k \log \left( \frac{1}{\varepsilon \eta} \right) \quad \text{and} \quad \eta := \frac{1}{\text{tr}(\mathbf{A})} \sum_{j>k} \sigma_j(\mathbf{A}).$$