Lecture 11: Concentration of sums of independent random variables

School of Mathematical Sciences, Xiamen University

1. Gaussian, or normal, distribution $\mathcal{N}(\mu, \sigma^2)$

- Standard normal random variable $X \sim \mathcal{N}(0, 1)$:

  (i) Tail:

$$\mathbb{P}\{|X| \geq t\} \leq 2\exp(-t^2/2), \quad \forall\, t \geq 0.$$

  (ii) Moment:

$$\|X\|_p := (\mathbb{E}|X|^p)^{1/p} = \mathcal{O}(\sqrt{p}) \quad \text{as } p \to \infty.$$

  (iii) Moment Generation function (MGF):

$$\mathbb{E}\exp(\lambda X) = \exp(\lambda^2/2) \quad \text{for all } \lambda \in \mathbb{R}.$$

  (iv) MGF of square:

$$\mathbb{E}\exp(cX^2) \leq 2 \quad \text{for some } c > 0.$$

- The sum of independent normal random variables is also normal.

2x. Sub-gaussian distributions

Theorem 1 (Sub-gaussian properties)

For a random variable X, the following properties are equivalent.

- *Tail:* $\mathbb{P}\{|X| \geq t\} \leq 2\exp(-t^2/K_1^2)$ *for all* $t \geq 0$.
- *Moment:* $\|X\|_p := (\mathbb{E}|X|^p)^{1/p} \leq K_2\sqrt{p}$ *for all* $p \geq 1$.
- *MGF of square:* $\mathbb{E}\exp(X^2/K_3^2) \leq 2$.

Moreover, if $\mathbb{E}X = 0$ *then these properties are also equivalent to the following one:*

- *MGF:* $\mathbb{E}\exp(\lambda X) \leq \exp(\lambda^2 K_4^2)$ *for all* $\lambda \in \mathbb{R}$.

Remark. The parameters $K_i > 0$ appearing in these properties can be different. However, they may differ from each other by at most an absolute constant factor. This means that there exists an absolute constant C such that property 1 implies property 2 with parameter $K_2 \leq CK_1$, and similarly for every other pair or properties.

Definition 2

Random variables that satisfy one of the first three properties (and thus all of them) in Theorem 1 are called *sub-gaussian*.

- The best K_3 is called the sub-gaussian norm of X, and is usually denoted $\|X\|_{\psi_2}$, that is

$$\|X\|_{\psi_2} := \inf\{t > 0 : \mathbb{E}\exp(X^2/t^2) \leq 2\}.$$

- Sub-gaussian random variable examples.

  (i) Normal random variables $X \sim \mathcal{N}(\mu, \sigma^2)$.

  (ii) Bernoulli random variable $X = 0, 1$ with probabilities $1/2$ each.

  (iii) More generally, any bounded random variable X.

- Not sub-gaussian random variable examples.

  Poisson, exponential, Pareto and Cauchy distributions.

Theorem 3 (Sums of sub-gaussians)

Let $X_1, \dots, X_N$ be independent, mean zero, sub-gaussian random variables. Then $\sum_{i=1}^N X_i$ is sub-gaussian, and

$$\left\| \sum_{i=1}^N X_i \right\|_{\psi_2}^2 \le C \sum_{i=1}^N \|X_i\|_{\psi_2}^2,$$

where C is an absolute constant.

Proof. Let us bound the MGF of the sum for any $\lambda \in \mathbb{R}$:

$$\mathbb{E} \exp\left(\lambda \sum_{i=1}^N X_i\right) = \prod_{i=1}^N \mathbb{E} \exp(\lambda X_i) \quad \text{(using independence)}$$

$$\le \prod_{i=1}^N \exp(C\lambda^2 \|X_i\|_{\psi_2}^2) \quad \text{(by last property in Theorem 1)}$$

$$= \exp(\lambda^2 K^2) \quad \text{where } K^2 := C \sum_{i=1}^N \|X_i\|_{\psi_2}^2. \quad \square$$

2.1. Hoeffding's inequality

- We rewrite Theorem 3 as a concentration inequality by using the first property in Theorem 1.

Theorem 4 (Hoeffding's inequality)

Let $X_1, \ldots, X_N$ be independent, mean zero, sub-gaussian random variables. Then, for every $t \geq 0$ we have

$$\mathbb{P}\left\{\left|\sum_{i=1}^{N} X_i\right| \geq t\right\} \leq 2\exp\left(-\frac{Ct^2}{\sum_{i=1}^{N}\|X_i\|_{\psi_2}^2}\right).$$

Remark 5

Hoeffding's inequality controls how far and with what probability a sum of independent random variables can deviate from its mean, which is zero.

3. Sub-exponential distributions

- The square X^2 of a normal random variable $X \sim \mathcal{N}(0,1)$ is not sub-gaussian.

Theorem 6 (Sub-exponential properties)

For a random variable X, the following properties are equivalent.

- *Tail:* $\mathbb{P}\{|X| \geq t\} \leq 2\exp(-t/K_1)$ *for all* $t \geq 0$.
- *Moment:* $\|X\|_p := (\mathbb{E}|X|^p)^{1/p} \leq pK_2$ *for all* $p \geq 1$.
- *MGF of square:* $\mathbb{E}\exp(|X|/K_3) \leq 2$.

Moreover, if $\mathbb{E}X = 0$ then these properties imply the following one:

- *MGF:* $\mathbb{E}\exp(\lambda X) \leq \exp(\lambda^2 K_4^2)$ *for* $|\lambda| \leq 1/K_4$.

Definition 7

Random variables that satisfy one of the first three properties (and thus all of them) in Theorem 6 are called *sub-exponential*.

- The best K_3 is called the sub-exponential norm of X, and is usually denoted $\|X\|_{\psi_1}$, that is

$$\|X\|_{\psi_1} := \inf\{t > 0 : \mathbb{E}\exp(|X|/t) \leq 2\}.$$

- All squares of sub-gaussian random variables are sub-exponential random variables. (Conversely?) We have

$$\|X^2\|_{\psi_1} = \|X\|_{\psi_2}^2.$$

3.1 Bernstein's inequality

Theorem 8 (Bernstein's inequality)

Let $X_1, \ldots, X_N$ be independent, mean zero, sub-exponential random variables. Then, for every $t \geq 0$ we have

$$\mathbb{P}\left\{\left|\sum_{i=1}^{N} X_i\right| \geqslant t\right\} \leqslant 2\exp\left[-C\min\left(\frac{t^2}{\sum_{i=1}^{N}\|X_i\|_{\psi_1}^2}, \frac{t}{\max_i \|X_i\|_{\psi_1}}\right)\right].$$

Proof. Choose $\lambda \geq 0$ and use Markov's inequality to get $(S = \sum_{i=1}^{N} X_i)$

$$\mathbb{P}\{S \geqslant t\} = \mathbb{P}\{\exp(\lambda S) \geqslant \exp(\lambda t)\} \leqslant \mathrm{e}^{-\lambda t}\mathbb{E}\exp(\lambda S).$$

Then by independence, we have

$$\mathbb{P}\{S \geqslant t\} \leq \mathrm{e}^{-\lambda t}\prod_{i=1}^{N}\mathbb{E}\exp(\lambda X_i).$$

If we choose λ small enough so that $0 < \lambda \leqslant \frac{C}{\max_i \|X_i\|_{\psi_1}}$, then by the last property in Theorem 6 we have

$$\mathbb{E}\exp\left(\lambda X_i\right) \leqslant \exp\left(C\lambda^2 \left\|X_i\right\|_{\psi_1}^2\right).$$

Hence,

$$\mathbb{P}\{S \geqslant t\} \leq \exp\left(-\lambda t + C\lambda^2\sigma^2\right), \quad \sigma^2 = \sum_{i=1}^{N}\left\|X_i\right\|_{\psi_1}^2.$$

The remaining part is left as an exercise. □

- Why does Bernstein's inequality have a mixture of two tails?

  (i) The sub-exponential tail should of course be there. Indeed, even if the entire sum consisted of a single term X_i, the best bound we could hope for would be of the form $\exp(-Ct/\|X_i\|_{\psi_1})$.

  (ii) The sub-gaussian term could be explained by the central limit theorem, which states that the sum should becomes approximately normal as the number of terms N increases to infinity.

Remark 9 (Bernstein's inequality for bounded random variables)

Suppose further the random variables X_i are uniformly bounded, which is a stronger assumption than being sub-gaussian. If $K > 0$ is such that $|X_i| \leq K$ almost surely for all i, then, for every $t \geq 0$, we have

$$\mathbb{P}\left\{\left|\sum_{i=1}^{N} X_i\right| \geqslant t\right\} \leqslant 2\exp\left(-\frac{t^2/2}{\sigma^2 + CKt}\right),$$

where $\sigma^2 = \sum_{i=1}^{N} \mathbb{E}X_i^2$ is the variance of the sum.

- Note that $\sigma^2 + CKt \leqslant 2\max\left(\sigma^2, CKt\right)$. So we can state the probability bound as

$$2\exp\left[-C\min\left(\frac{t^2}{\sigma^2}, \frac{t}{K}\right)\right].$$

  Just as before, here we also have a mixture of two tails, sub-gaussian and sub-exponential.

  The sub-gaussian tail is a bit *sharper* than in Theorem 8, since it depends on the variances rather than sub-exponential norms of X_i.

  The sub-exponential tail, on the other hand, is *weaker*, since it depends on the sup-norms rather than the sub-exponential norms of X_i.

- **More on concentration.** 

4. Sub-gaussian random vectors

- Definition. Consider a random vector $\boldsymbol{X}$ taking values in $\mathbb{R}^n$. We call $\boldsymbol{X}$ a sub-gaussian random vector if all one-dimensional marginals of $\boldsymbol{X}$, i.e., the random variables $\langle \boldsymbol{X}, \mathbf{x} \rangle$ for all $\mathbf{x} \in \mathbb{R}^n$, are sub-gaussian.
- The sub-gaussian norm of $\boldsymbol{X}$ is defined as

$$\|\boldsymbol{X}\|_{\psi_2} := \sup_{\mathbf{x} \in \mathbb{R}^n, \|\mathbf{x}\|_2 = 1} \|\langle \boldsymbol{X}, \mathbf{x} \rangle\|_{\psi_2}.$$

- Sub-gaussian random vector examples

  (i) The standard normal distribution $\mathcal{N}(\mathbf{0}, \mathbf{I}_n)$ (why?)

  (ii) The uniform distribution on the centered Euclidean sphere of radius $\sqrt{n}$

  (iii) The uniform distribution on the cube $\{-1, 1\}^n$

  (iv) A random vector $\boldsymbol{X} = (X_1, \cdots, X_n)$ with independent and sub-gaussian coordinates is sub-gaussian, $\|\boldsymbol{X}\|_{\psi_2} \leqslant C \max_i \|X_i\|_{\psi_2}$.

5. Johnson-Lindenstrauss Lemma

- Concentration inequalities like Hoeffding's and Bernstein's are successfully used in the analysis of algorithms.
- Let us give one example for the problem of dimension reduction. Suppose we have some data that is represented as a set of N points in $\mathbb{R}^n$. We would like to compress the data by representing it in a lower dimensional space $\mathbb{R}^m$ instead of $\mathbb{R}^n$ with $m \ll n$. By how much can we reduce the dimension without loosing the important features of the data?
- The basic result in this direction is the Johnson-Lindenstrauss Lemma. It states that a remarkably simple dimension reduction method works - a random linear map from $\mathbb{R}^n$ to $\mathbb{R}^m$ with $m \sim \log N$. The logarithmic function grows very slowly, so we can usually reduce the dimension dramatically.
- What exactly is a random linear map? We consider an $m \times n$ matrix $\boldsymbol{A}$ whose rows are independent, mean zero, *isotropic* ($\mathbb{E}\left[\boldsymbol{X}\boldsymbol{X}^\top\right] = \mathbf{I}_n$) and sub-gaussian random vectors in $\mathbb{R}^n$.

Theorem 10 (Johnson-Lindenstrauss Lemma)

Let $\mathcal{X}$ be a set of N points in $\mathbb{R}^n$ and $\varepsilon \in (0,1)$. Consider an $m \times n$ matrix $\boldsymbol{A}$ whose rows are independent, mean zero, isotropic and sub-gaussian random vectors $\boldsymbol{X}_i$ in $\mathbb{R}^n$. Rescale $\boldsymbol{A}$ by defining the "Gaussian random projection"

$$\boldsymbol{P} := \boldsymbol{A}/\sqrt{m}.$$

Assume that

$$m \geq C\varepsilon^{-2} \log N,$$

where C is an appropriately large constant that depends only on the sub-gaussian norms of the vectors $\boldsymbol{X}_i$. Then, with high probability (say, 0.99), the map $\boldsymbol{P}$ preserves the distances between all points in $\mathcal{X}$ with error ε, that is for all $\mathbf{x}, \mathbf{y} \in \mathcal{X}$,

$$(1-\varepsilon)\|\mathbf{x}-\mathbf{y}\|_2 \leqslant \|\boldsymbol{P}\mathbf{x}-\boldsymbol{P}\mathbf{y}\|_2 \leqslant (1+\varepsilon)\|\mathbf{x}-\mathbf{y}\|_2.$$

Examples of $\boldsymbol{A}$: Gaussian random matrix, `randn(m,n)` in MATLAB; an $m \times n$ matrix with independent Rademacher entries (± 1 with equal probabilities).

Proof. By linearity of $\boldsymbol{P}$, $1-\varepsilon \geq (1-\varepsilon)^2$, and $1+\varepsilon \leq (1+\varepsilon)^2$, it is sufficient to prove that (with high probability)

$$1-\varepsilon \leq \|\boldsymbol{P}\mathbf{z}\|_2^2 \leq 1+\varepsilon \quad \text{for all } \mathbf{z} \in \mathcal{T}$$

where

$$\mathcal{T} := \left\{ \frac{\mathbf{x}-\mathbf{y}}{\|\mathbf{x}-\mathbf{y}\|_2} : \mathbf{x}, \mathbf{y} \in \mathcal{X} \text{ and } \mathbf{x} \neq \mathbf{y} \right\}.$$

By $\boldsymbol{P}\mathbf{z} = \boldsymbol{A}\mathbf{z}/\sqrt{m}$, it is enough to show that (with high probability)

$$\left| \frac{1}{m} \sum_{i=1}^{m} \langle \boldsymbol{X}_i, \mathbf{z} \rangle^2 - 1 \right| \leq \varepsilon \quad \text{for all } \mathbf{z} \in \mathcal{T}.$$

We can prove this inequality by combining concentration and a union bound.

In order to use concentration, we first fix $\mathbf{z} \in \mathcal{T}$. By assumption, the random variables $\langle \boldsymbol{X}_i, \mathbf{z} \rangle^2 - 1$ are independent; they have zero mean (why? Exercise), and they are sub-exponential (why? Exercise). Then Bernstein's inequality gives (why? Exercise)

$$\mathbb{P}\left\{ \left| \frac{1}{m} \sum_{i=1}^{m} \langle \boldsymbol{X}_i, \mathbf{z} \rangle^2 - 1 \right| > \varepsilon \right\} \leqslant 2 \exp\left(-c\varepsilon^2 m\right).$$

Finally, we can unfix $\mathbf{z}$ by taking a union bound over all possible $\mathbf{z} \in \mathcal{T}$:

$$\begin{aligned} \mathbb{P}\left\{ \max_{\mathbf{z} \in \mathcal{T}} \left| \frac{1}{m} \sum_{i=1}^{m} \langle \boldsymbol{X}_i, \mathbf{z} \rangle^2 - 1 \right| > \varepsilon \right\} &\leqslant \sum_{\mathbf{z} \in \mathcal{T}} \mathbb{P}\left\{ \left| \frac{1}{m} \sum_{i=1}^{m} \langle \boldsymbol{X}_i, \mathbf{z} \rangle^2 - 1 \right| > \varepsilon \right\} \\ &\leqslant |\mathcal{T}| \cdot 2 \exp\left(-c\varepsilon^2 m\right). \end{aligned}$$

By definition of $\mathcal{T}$, we have $|\mathcal{T}| \leq N^2$. So, if we choose $m \geq C\varepsilon^{-2} \log N$ with appropriately large constant C, we can make

$$|\mathcal{T}| \cdot 2 \exp\left(-c\varepsilon^2 m\right) \leq 0.01.$$

The proof is complete. □